

Final_proj_3

April 13, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[67]: hosp = pd.read_csv('Hospitalisation details.csv')
```

```
[68]: medic = pd.read_csv('Medical Examinations.csv')
```

```
[4]: names = pd.read_excel('Names.xlsx')
```

```
[5]: hosp.head()
```

```
[5]:
```

	Customer ID	year	month	date	children	charges	Hospital tier	City tier	\
0	Id2335	1992	Jul	9	0	563.84	tier - 2	tier - 3	
1	Id2334	1992	Nov	30	0	570.62	tier - 2	tier - 1	
2	Id2333	1993	Jun	30	0	600.00	tier - 2	tier - 1	
3	Id2332	1992	Sep	13	0	604.54	tier - 3	tier - 3	
4	Id2331	1998	Jul	27	0	637.26	tier - 3	tier - 3	

```
State ID
0    R1013
1    R1013
2    R1013
3    R1013
4    R1013
```

```
[6]: final_data=pd.merge(hosp, medic,how='inner', on = 'Customer ID')
```

```
[7]: final_data.head()
```

```
[7]:
```

	Customer ID	year	month	date	children	charges	Hospital tier	City tier	\
0	Id2335	1992	Jul	9	0	563.84	tier - 2	tier - 3	
1	Id2334	1992	Nov	30	0	570.62	tier - 2	tier - 1	
2	Id2333	1993	Jun	30	0	600.00	tier - 2	tier - 1	
3	Id2332	1992	Sep	13	0	604.54	tier - 3	tier - 3	
4	Id2331	1998	Jul	27	0	637.26	tier - 3	tier - 3	

	State ID	BMI	HBA1C	Heart Issues	Any Transplants	Cancer history	\
0	R1013	17.58	4.51	No	No	No	
1	R1013	17.60	4.39	No	No	No	
2	R1013	16.47	6.35	No	No	Yes	
3	R1013	17.70	6.28	No	No	No	
4	R1013	22.34	5.57	No	No	No	

	NumberOfMajorSurgeries	smoker
0	1	No
1	1	No
2	1	No
3	1	No
4	1	No

```
[8]: final_data = final_data.merge(names,on='Customer ID')
```

```
[9]: final_data.shape
```

```
[9]: (2335, 17)
```

```
[10]: final_data.columns = final_data.columns.str.lower()
final_data.columns = final_data.columns.str.replace(' ', '_')
final_data.columns
```

```
[10]: Index(['customer_id', 'year', 'month', 'date', 'children', 'charges',
            'hospital_tier', 'city_tier', 'state_id', 'bmi', 'hba1c',
            'heart_issues', 'any_transplants', 'cancer_history',
            'numberofmajorsurgeries', 'smoker', 'name'],
            dtype='object')
```

```
[11]: (final_data=='?').sum()
```

```
[11]: customer_id      0
      year            2
      month           3
      date            0
      children         0
      charges          0
      hospital_tier    1
      city_tier        1
      state_id         2
      bmi             0
      hba1c            0
      heart_issues     0
      any_transplants  0
      cancer_history   0
```

```
numberofmajorsurgeries    0
smoker                    2
name                      0
dtype: int64
```

```
[12]: miss_perc = (final_data == '?').sum(axis = 1)/final_data.shape[1]
```

```
[13]: miss_perc[miss_perc > 0]
```

```
[13]: 11      0.058824
      13      0.058824
      17      0.117647
      542     0.058824
      1046    0.058824
      1049    0.058824
      1700    0.058824
      1775    0.058824
      2165    0.058824
      2332    0.058824
dtype: float64
```

```
[14]: miss_perc_colm = (final_data == '?').sum(axis = 0)/final_data.shape[0]
```

```
[15]: miss_perc_colm.sort_values(ascending=False)
```

```
[15]: month                0.001285
      state_id             0.000857
      year                 0.000857
      smoker               0.000857
      city_tier             0.000428
      hospital_tier         0.000428
      date                 0.000000
      children              0.000000
      charges               0.000000
      name                 0.000000
      bmi                  0.000000
      hba1c                 0.000000
      heart_issues          0.000000
      any_transplants        0.000000
      cancer_history         0.000000
      numberofmajorsurgeries 0.000000
      customer_id           0.000000
dtype: float64
```

```
[16]: final_noq = final_data.drop(index = miss_perc[miss_perc>0].index)
```

```
[17]: final_noq.shape
```

```
[17]: (2325, 17)
```

```
[18]: final_noq[['city_tier', 'hospital_tier']]
```

```
[18]:   city_tier hospital_tier
0      tier - 3      tier - 2
1      tier - 1      tier - 2
2      tier - 1      tier - 2
3      tier - 3      tier - 3
4      tier - 3      tier - 3
...
2329   tier - 3      tier - 1
2330   tier - 2      tier - 1
2331   tier - 3      tier - 1
2333   tier - 3      tier - 2
2334   tier - 3      tier - 1
```

```
[2325 rows x 2 columns]
```

```
[19]: from sklearn.preprocessing import OrdinalEncoder
ordinal = OrdinalEncoder(categories=[['tier - 3','tier - 2','tier - 1'],[
    'tier - 3', 'tier - 2', 'tier - 1']])
final_noq[['city_tier_ord','hospital_tier_ord']] = ordinal.
    fit_transform(final_noq[['city_tier','hospital_tier']])
```

```
[20]: pd.crosstab(final_noq['city_tier_ord'],final_noq['city_tier'])
```

```
[20]: city_tier      tier - 1  tier - 2  tier - 3
city_tier_ord
0.0              0         0       789
1.0              0       807         0
2.0             729         0         0
```

```
[21]: pd.crosstab(final_noq['hospital_tier_ord'],final_noq['hospital_tier'])
```

```
[21]: hospital_tier      tier - 1  tier - 2  tier - 3
hospital_tier_ord
0.0              0         0       691
1.0              0      1334         0
2.0             300         0         0
```

```
[22]: final_noq.head()
```

```
[22]:   customer_id  year month  date  children  charges hospital_tier city_tier \
0      Id2335  1992   Jul    9         0   563.84      tier - 2  tier - 3
1      Id2334  1992  Nov   30         0   570.62      tier - 2  tier - 1
2      Id2333  1993   Jun   30         0   600.00      tier - 2  tier - 1
```

3	Id2332	1992	Sep	13	0	604.54	tier - 3	tier - 3
4	Id2331	1998	Jul	27	0	637.26	tier - 3	tier - 3

	state_id	bmi	hba1c	heart_issues	any_transplants	cancer_history	\
0	R1013	17.58	4.51	No	No	No	
1	R1013	17.60	4.39	No	No	No	
2	R1013	16.47	6.35	No	No	Yes	
3	R1013	17.70	6.28	No	No	No	
4	R1013	22.34	5.57	No	No	No	

	numberofmajorsurgeries	smoker	name	\
0	1	No	German, Mr. Aaron K	
1	1	No	Rosendahl, Mr. Evan P	
2	1	No	Albano, Ms. Julie	
3	1	No	Riveros Gonzalez, Mr. Juan D. Sr.	
4	1	No	Brietzke, Mr. Jordan	

	city_tier_ord	hospital_tier_ord
0	0.0	1.0
1	2.0	1.0
2	2.0	1.0
3	0.0	0.0
4	0.0	0.0

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[23]: final_noq.head()
```

```
[23]:
```

	customer_id	year	month	date	children	charges	hospital_tier	city_tier	\
0	Id2335	1992	Jul	9	0	563.84	tier - 2	tier - 3	
1	Id2334	1992	Nov	30	0	570.62	tier - 2	tier - 1	
2	Id2333	1993	Jun	30	0	600.00	tier - 2	tier - 1	
3	Id2332	1992	Sep	13	0	604.54	tier - 3	tier - 3	
4	Id2331	1998	Jul	27	0	637.26	tier - 3	tier - 3	

	state_id	bmi	hba1c	heart_issues	any_transplants	cancer_history	\
0	R1013	17.58	4.51	No	No	No	
1	R1013	17.60	4.39	No	No	No	
2	R1013	16.47	6.35	No	No	Yes	
3	R1013	17.70	6.28	No	No	No	
4	R1013	22.34	5.57	No	No	No	

	numberofmajorsurgeries	smoker	name \
0	1	No	German, Mr. Aaron K
1	1	No	Rosendahl, Mr. Evan P
2	1	No	Albano, Ms. Julie
3	1	No	Riveros Gonzalez, Mr. Juan D. Sr.
4	1	No	Brietzke, Mr. Jordan

	city_tier_ord	hospital_tier_ord
0	0.0	1.0
1	2.0	1.0
2	2.0	1.0
3	0.0	0.0
4	0.0	0.0

```
[24]: vc = final_noq.state_id.value_counts() # frequency of each category
      vc[:3].index
```

```
[24]: Index(['R1013', 'R1011', 'R1012'], dtype='object')
```

```
[25]: for i in vc[:3].index:
      var_name = 'state_id_' + i
      print(var_name)
      final_noq[var_name] = 0
      final_noq.loc[final_noq.state_id == i, var_name] = 1
```

```
state_id_R1013
state_id_R1011
state_id_R1012
```

```
[26]: final_noq.state_id.value_counts()
```

```
[26]: R1013    609
      R1011    574
      R1012    572
      R1024    159
      R1026     84
      R1021     70
      R1016     64
      R1025     40
      R1023     38
      R1017     36
      R1019     26
      R1022     14
      R1014     13
      R1015     11
      R1018      9
```

```
R1020      6
Name: state_id, dtype: int64
```

```
[27]: final_noq['state_id_R1013'].value_counts()
```

```
[27]: 0    1716
      1     609
      Name: state_id_R1013, dtype: int64
```

```
[28]: final_noq.numberofmajorsurgeries.unique()
```

```
[28]: array(['1', 'No major surgery', '2', '3'], dtype=object)
```

```
[29]: final_noq.loc[final_noq.numberofmajorsurgeries == 'No major_
      ↪surgery', 'numberofmajorsurgeries' ] = 0
```

```
[30]: final_noq.numberofmajorsurgeries = final_noq.numberofmajorsurgeries.astype(int)
```

```
[31]: final_noq.year = final_noq.year.astype(int)
```

```
[32]: final_noq['age'] = 2023- final_noq.year
```

```
[33]: final_noq['title'] = final_noq.name.str.split('[,.]').str[1].str.strip()
```

```
[34]: final_noq.title.value_counts()
```

```
[34]: Mr      1160
      Ms      1023
      Mrs     142
      Name: title, dtype: int64
```

```
[35]: final_noq['gender'] = 'female'
      final_noq.loc[final_noq.title == 'Mr', 'gender' ] = 'male'
```

```
[36]: final_noq.loc[final_noq.title == 'Mrs']
```

```
[36]:
```

	customer_id	year	month	date	children	charges	hospital_tier	\
24	Id2311	2001	Aug	19	0	964.71	tier - 3	
172	Id2163	2004	Dec	27	0	1863.45	tier - 3	
197	Id2138	2004	Jun	12	0	2094.10	tier - 3	
328	Id2007	1993	Sep	25	0	3162.02	tier - 2	
348	Id1987	2003	Dec	5	0	3300.70	tier - 2	
...		
1790	Id545	1963	Jul	4	0	18208.34	tier - 1	
1808	Id527	1963	Dec	6	0	18883.33	tier - 1	
1811	Id524	1963	Oct	20	0	18954.56	tier - 1	
1839	Id496	1966	Aug	10	0	19995.29	tier - 1	

1848		Id487	1962	Jul	2	0	20354.50	tier - 3
------	--	-------	------	-----	---	---	----------	----------

	city_tier	state_id	bmi	...	smoker	name	\
24	tier - 2	R1013	25.19	...	No	Keys, Mrs. Kathleen	
172	tier - 1	R1025	27.06	...	No	Stanislav, Mrs. Grace H	
197	tier - 2	R1025	27.74	...	No	Padula, Mrs. Lauren	
328	tier - 3	R1013	25.61	...	No	Martin, Mrs. Kristen M	
348	tier - 2	R1025	30.54	...	No	Mendez-Karr, Mrs. Cynthia	
...	
1790	tier - 2	R1026	44.20	...	No	Shigezumi, Mrs. Teiko	
1808	tier - 1	R1026	46.19	...	No	Hughey, Mrs. Ashley E	
1811	tier - 1	R1026	46.40	...	No	Rogers, Mrs. Anita L.	
1839	tier - 3	R1026	51.74	...	No	Oehlke, Mrs. Jessica	
1848	tier - 2	R1026	49.77	...	No	Argall, Mrs. Tara R	

	city_tier_ord	hospital_tier_ord	state_id_R1013	state_id_R1011	\
24	1.0	0.0	1	0	
172	2.0	0.0	0	0	
197	1.0	0.0	0	0	
328	0.0	1.0	1	0	
348	1.0	1.0	0	0	
...	
1790	1.0	2.0	0	0	
1808	2.0	2.0	0	0	
1811	2.0	2.0	0	0	
1839	0.0	2.0	0	0	
1848	1.0	0.0	0	0	


	state_id_R1012	age	title	gender
24	0	22	Mrs	female
172	0	19	Mrs	female
197	0	19	Mrs	female
328	0	30	Mrs	female
348	0	20	Mrs	female
...
1790	0	60	Mrs	female
1808	0	60	Mrs	female
1811	0	60	Mrs	female
1839	0	57	Mrs	female
1848	0	61	Mrs	female

[142 rows x 25 columns]

```
[37]: plt.figure(figsize = (25,10))
plt.hist(final_noq.charges, bins = 100)
plt.xlabel('Charges')
plt.title('Histogram')
```




```
plt.show()
```




```
[38]: plt.figure(figsize = (25,10))
sns.boxplot(final_noq.charges)
plt.xlabel('Charges')
plt.title('Box-Plot')
plt.show()
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


FutureWarning




```
[39]: plt.figure(figsize =(25,10))
sns.boxplot( x='charges',y='gender', data= final_noq)
plt.show()
```



```
[40]: plt.figure(figsize =(25,10))
sns.boxplot( x='charges',y='city_tier', data= final_noq)
plt.show()
```




```
[41]: plt.figure(figsize =(25,10))
sns.boxplot( x='charges',y='hospital_tier', data= final_noq)
plt.show()
```



```
[42]: median = final_noq.groupby('hospital_tier')[['charges']].median().reset_index()
```

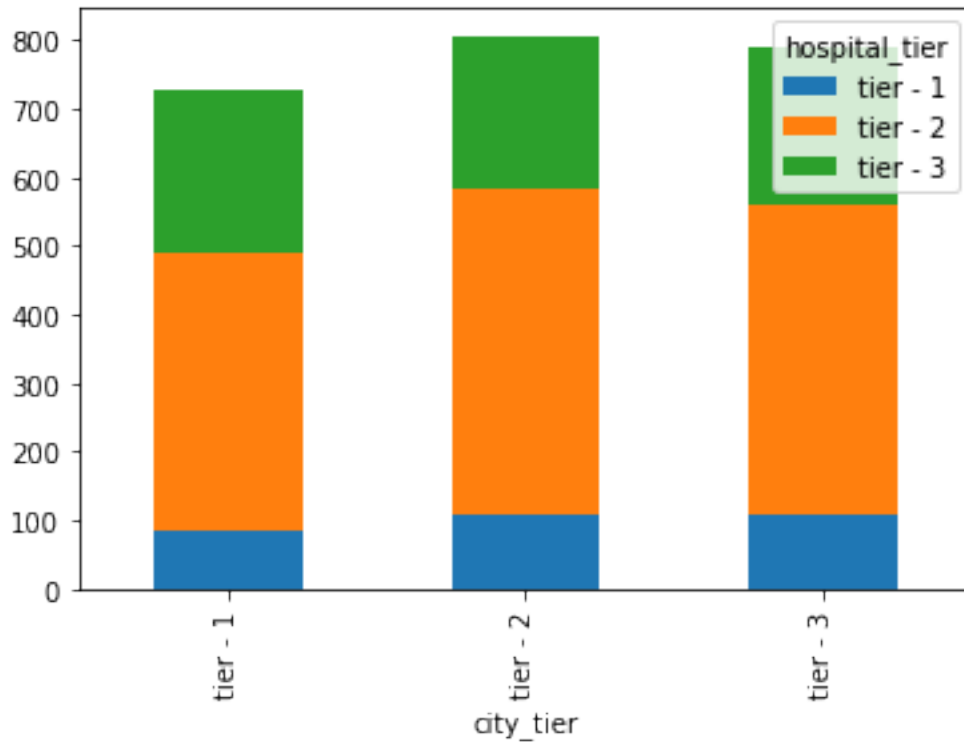
```
[43]: import plotly.express as px
fig = px.line_polar(median, r='charges', theta='hospital_tier', line_close=True)
fig.show()
```



```
[44]: pd.crosstab(final_noq.city_tier, final_noq.hospital_tier)
```

```
[44]: hospital_tier  tier - 1  tier - 2  tier - 3
city_tier
tier - 1           85       403       241
tier - 2          106       479       222
tier - 3          109       452       228
```

```
[45]: pd.crosstab(final_noq.city_tier, final_noq.hospital_tier).plot.bar(stacked=True)
plt.show()
```



```
[46]: from statsmodels.formula.api import ols
import statsmodels.api as sm

mod = ols('charges ~ hospital_tier', data = final_noq).fit()
res = sm.stats.anova_lm(mod)
```

```
[47]: from statsmodels.stats.multicomp import pairwise_tukeyhsd

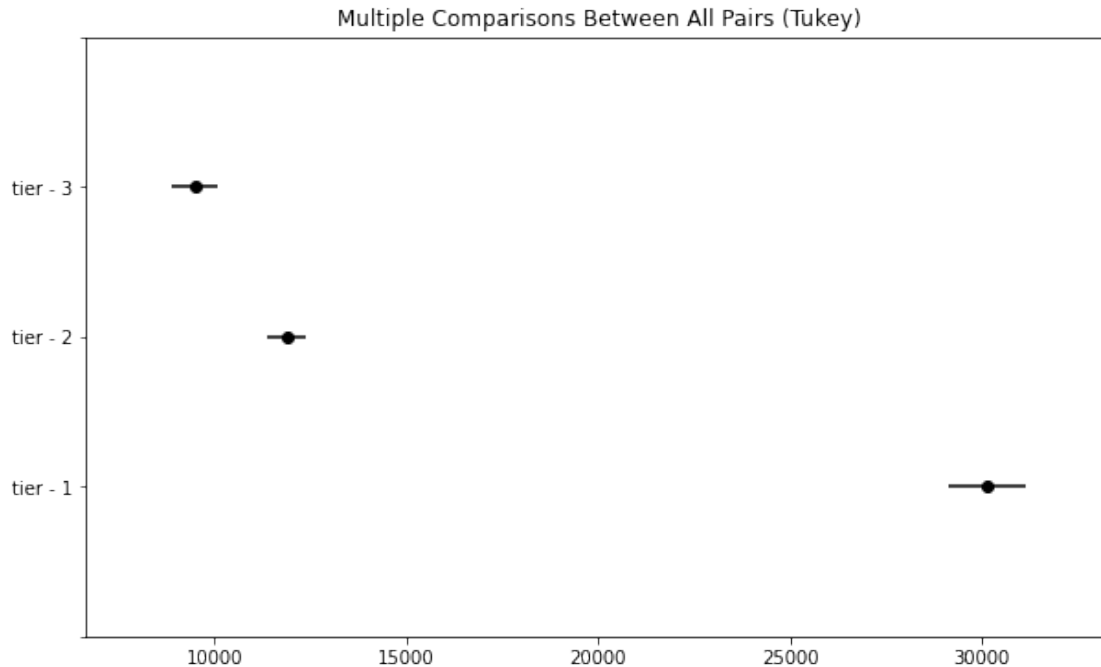
res_tukey = pairwise_tukeyhsd(final_noq.charges, final_noq.hospital_tier)
res_tukey.summary()
```

```
[47]: <class 'statsmodels.iolib.table.SimpleTable'>
```

```
[48]: res_tukey.plot_simultaneous()
plt.show()
```

```
/usr/local/lib/python3.7/site-
packages/statsmodels/sandbox/stats/multicomp.py:775: UserWarning:
```

```
FixedFormatter should only be used together with FixedLocator
```



```
[49]: import scipy.stats as stats
```

```
sample1 = final_noq.loc[final_noq.smoker == 'yes', 'charges']
sample2 = final_noq.loc[final_noq.smoker != 'yes', 'charges']
stats.ttest_ind(sample1, sample2)
```

```
[49]: Ttest_indResult(statistic=74.15560699695726, pvalue=0.0)
```

```
[50]: mod = ols('charges ~ city_tier', data = final_noq).fit()
res = sm.stats.anova_lm(mod)
res
```

```
[50]:
```

	df	sum_sq	mean_sq	F	PR(>F)
city_tier	2.0	4.092192e+08	2.046096e+08	1.454356	0.233763
Residual	2322.0	3.266763e+11	1.406874e+08	NaN	NaN

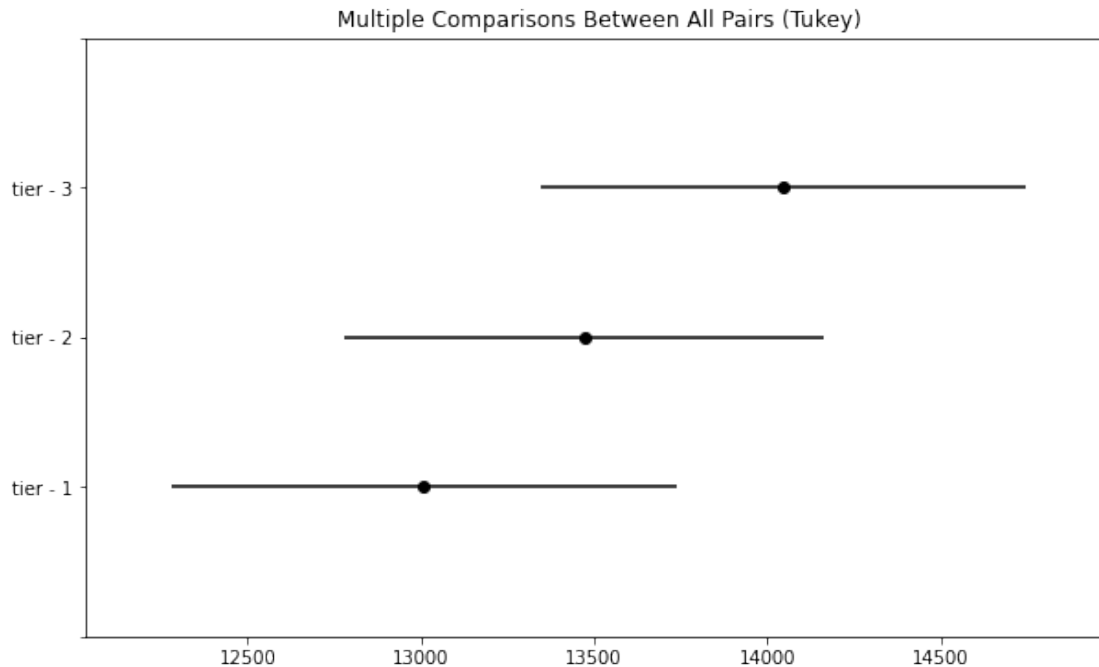
```
[51]: res_tukey = pairwise_tukeyhsd(final_noq.charges, final_noq.city_tier)
res_tukey.summary()
```

```
[51]: <class 'statsmodels.iolib.table.SimpleTable'>
```

```
[52]: res_tukey.plot_simultaneous()
plt.show()
```

```
/usr/local/lib/python3.7/site-
```

```
packages/statsmodels/sandbox/stats/multicomp.py:775: UserWarning:
FixedFormatter should only be used together with FixedLocator
```



```
[53]: observed_table = pd.crosstab(final_noq.smoker, final_data.heart_issues)
```

```
[54]: chi, p, df, expected = stats.chi2_contingency(observed_table)
```

```
[55]: chi, p, df, expected
```

```
[55]: (0.08588150449910657,
0.7694797581780767,
1,
array([[1111.30967742, 727.69032258],
[ 293.69032258, 192.30967742]]))
```

```
[56]: data = final_noq.drop(columns = ['name', 'year', 'month', '
↳ 'date', 'hospital_tier',
'city_tier', 'state_id', 'title'])
```

```
[57]: corr_plot = data.select_dtypes(exclude='object').corr()
ma = np.ones_like(corr_plot)
ma[np.tril_indices_from(ma)] = 0
```

```
[58]: data_2 = pd.get_dummies(data, drop_first=True)
data_2.reset_index(drop=True, inplace = True)
```

```
[59]: model_data = data_2.drop(columns = 'charges')
model_data.head()
model_data['charges'] = data_2.charges
model_data.head()
```

```
[59]:
```

	children	bmi	hba1c	numberofmajorsurgeries	city_tier_ord	\
0	0	17.58	4.51	1	0.0	
1	0	17.60	4.39	1	2.0	
2	0	16.47	6.35	1	2.0	
3	0	17.70	6.28	1	0.0	
4	0	22.34	5.57	1	0.0	

	hospital_tier_ord	state_id_R1013	state_id_R1011	state_id_R1012	age	\
0	1.0	1	0	0	31	
1	1.0	1	0	0	31	
2	1.0	1	0	0	30	
3	0.0	1	0	0	31	
4	0.0	1	0	0	25	

	...	customer_id_Id996	customer_id_Id997	customer_id_Id998	\
0	...	0	0	0	
1	...	0	0	0	
2	...	0	0	0	
3	...	0	0	0	
4	...	0	0	0	

	customer_id_Id999	heart_issues_yes	any_transplants_yes	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	cancer_history_Yes	smoker_yes	gender_male	charges
0	0	0	1	563.84
1	0	0	1	570.62
2	1	0	0	600.00
3	0	0	1	604.54
4	0	0	1	637.26


```
[5 rows x 2340 columns]
```

```
[60]: model_data.columns = model_data.columns.str.lower()
```

```
[61]: # converting y to categorical for stratified k fold
y_class = pd.cut(model_data.charges, bins = 4, labels= [1,2,3,4])
X = model_data.drop(columns = 'charges')

[62]: from sklearn.model_selection import KFold, StratifiedKFold, RandomizedSearchCV,
      ↪ train_test_split
from sklearn.ensemble import RandomForestRegressor

from sklearn.linear_model import SGDRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error as mse, r2_score
from sklearn.preprocessing import StandardScaler
import xgboost as xgb
from sklearn.pipeline import Pipeline
from xgboost import XGBRegressor
folds = StratifiedKFold(n_splits=5, shuffle = True, random_state=12)
sgd_rmse_train = {}
sgd_rmse_test = {}
i= 1

[63]: for train_index, test_index in folds.split(X,y_class):

    train, test = model_data.loc[train_index,], model_data.loc[test_index,]

    # standardization :
    sc = StandardScaler()
    sc.fit(train)
    train_std = sc.transform(train)
    test_std = sc.transform(test)

    x_train , x_test = train_std[:, :-1], test_std[:, :-1]
    y_train, y_test = train_std[:, -1], test_std[:, -1]
    sgd = SGDRegressor(max_iter=100)

[64]: space = dict()
space['penalty'] = ['l1', 'l2', 'elasticnet']
space['l1_ratio'] = [0,.1,.2,.8,1]
space['alpha'] = [1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 1, 10, 100, 1000,10000]
space['learning_rate'] = ['constant', 'adaptive']
space['eta0']=[1e-5, 1e-4, 1e-3, 1e-2, 1e-1 , 2e-1, 3e-1, 5e-1, 8e-1, 4e-1,
↪ 8e-1, 1, 10, 100]

[65]: search = RandomizedSearchCV(sgd, space,
                                cv=5, refit=True, return_train_score = True,
```



```
random_state = 12)
```

```
[ ]:
```

```
[66]: earch = RandomizedSearchCV(xgb, space, scoring='neg_root_mean_squared_error',
                                cv=fold_inner, refit=True, return_train_score=
                                ↪= True,
                                random_state = 12)
```

```
↪-----
```

```
NameError                                Traceback (most recent call↪
↪last)
```

```
<ipython-input-66-8055994c2f9c> in <module>
      1 earch = RandomizedSearchCV(xgb, space,↪
↪scoring='neg_root_mean_squared_error',
----> 2                                cv=fold_inner, refit=True,↪
↪return_train_score = True,
      3                                random_state = 12)
```

```
NameError: name 'fold_inner' is not defined
```

```
[ ]: result = search.fit(x_train, y_train)
```

```
train_pred = result.predict(x_train)
test_pred = result.predict(x_test)
```

```
[ ]: sgd_rmse_train.update({'Fold{ }'.format(i): round(mse(y_true = y_train, y_pred =↪
↪train_pred, squared = False),3)})
sgd_rmse_test.update({'Fold{ }'.format(i): round(mse(y_true = y_test, y_pred =↪
↪test_pred, squared = False),3)})
i += 1
```

```
[ ]: train, test = train_test_split(model_data, test_size = 0.25, random_state = 12)
```

```
# standardization :
sc = StandardScaler()
sc.fit(train)
trn_std = sc.transform(train)
tst_std = sc.transform(test)
```

```

# getting X and y :
x_train, x_test = trn_std[:, :-1], tst_std[:, :-1]
y_train, y_test = trn_std[:, -1], tst_std[:, -1]

# sgd regression with hyperparameter tuning :
rf = RandomForestRegressor(random_state = 12)

# define search space
space = dict()
space['n_estimators'] = [10, 100, 500]
space['max_features'] = [2, 3, 4, 5, 6]

# define search
search = RandomizedSearchCV(rf, space, scoring='neg_root_mean_squared_error',
                             cv=5, refit=True, return_train_score = True,
                             random_state = 12, n_iter = 3
                             )

# execute search
result = search.fit(x_train, y_train)

```

```
[ ]: importance = pd.Series(result.best_estimator_.feature_importances_, index =
    ↳ train.drop(columns= 'charges').columns)
```

```
[ ]: impvars = importance.sort_values(ascending=False)[:6]
```

```
[ ]: pred_data = pd.DataFrame({'Name' : ['Christopher, Ms. Jayna'],
    'DOB' : ['12/28/1988'],
    'city_tier' : ['tier - 1'], 'children' : [ 2],
    'HbA1c' : [5.8],
    'smoker_yes' : [1],
    'heart_issues_yes' : [0],
    'any_transplants_yes' : [0],
    'numberofmajorsurgeries' : [ 0],
    'cancer_history_yes' : [1],
    'hospital_tier' : ['tier - 1'],
    'bmi' : [85/(1.70 **2)],
    'state_id_R1011' : [1]
    })
```

```
[ ]: pred_data
```

```
[ ]:
```